


```
In [37]: train.head(10).min(axis=1)
```

```
Out[37]: 0      0
          1      0
          2      0
          3      0
          4      0
          5      0
          6      0
          7      0
          8      0
          9      0
          dtype: int64
```


```
In [36]: train.head(10).max(axis=1)
```

```
Out[36]: 0      255
          1      255
          2      255
          3      255
          4      254
          5      255
          6      255
          7      255
          8      255
          9      255
          dtype: int64
```

As shown above, the image values range from 0 to 255 (black to white).

Next we can preview the images as 28x28 arrays.

```
In [30]: fig, ax = plt.subplots(10,10)
for i in range(10):
    for j in range(10):
        ax[i,j].imshow(np.reshape(np.array(train.iloc[10*i+j,1:]),(28,28)))
        ax[i,j].set_xticks([])
        ax[i,j].set_yticks([])
        ax[i,j].set_xlabel(train.iloc[10*i+j,0])
fig.set_size_inches((8,8))
fig.tight_layout()
```



The numbers in the images all are fairly centered and scaled to about the same size, and they are mostly black and white (not much grey), so a simple neural network should work for this data set.

Before modeling, let's normalize all the images.

```
In [39]: train.iloc[:,1:]=train.iloc[:,1:]/255
test = test/255
```

Modeling

To avoid overfitting, we should perform cross validation to select model hyperparameters. We will do a 5-fold CV.

```
In [56]: def CVsplit(i):  
    if i==1:  
        return (train.iloc[:8400,:],train.iloc[8400:,:])  
    elif i==5:  
        return (train.iloc[33600:,:],train.iloc[:33600,:])  
    else:  
        return (train.iloc[8400*(i-1):8400*i,:], train.iloc[:8400*(i-1),:].append(tr
```

This function will perform the CV given specified hyperparameters.

```
In [78]: def evalmodel(hidden, activation1, activation2, optimizer1, epochs1):  
    meanacc = 0  
    for i in range(1,6):  
        CVsets = CVsplit(i)  
        model = keras.Sequential([keras.layers.Dense(hidden, activation=activation1)  
            keras.layers.Dense(10, activation = activation2)])  
        model.compile(optimizer=optimizer1, loss='sparse_categorical_crossentropy',  
            model.fit(np.array(CVsets[1].iloc[:,1:]), np.array(CVsets[1].label), epochs=  
            resultloss, resultacc = model.evaluate(np.array(CVsets[0].iloc[:,1:]), np.ar  
            meanacc += resultacc  
    meanacc /=5  
    print('Hidden size',hidden, ', Activation 1:', str(activation1).split(' ')[1],',
```

Hyperparameter optimization is performed below. First the activation functions are optimized.

```
In [85]: for j in [tf.nn.relu,tf.nn.softmax]:
        for k in [tf.nn.relu,tf.nn.softmax]:
            evalmodel(125,j,k,'adam',5)
```

```
8400/8400 [=====] - 2s 226us/sample - loss: 2.2927 - acc:
0.1032
8400/8400 [=====] - 2s 224us/sample - loss: 2.3012 - acc:
0.0921
8400/8400 [=====] - 2s 246us/sample - loss: 2.3020 - acc:
0.1021
8400/8400 [=====] - 2s 247us/sample - loss: 2.2850 - acc:
0.1055
8400/8400 [=====] - 2s 257us/sample - loss: 2.3026 - acc:
0.1014
Hidden size 125 , Activation 1: relu , Activation 2: relu , Optimizer: adam , 5 epo
chs. Accuracy: 0.10088095217943191
8400/8400 [=====] - 2s 229us/sample - loss: 0.1034 - acc:
0.9696
8400/8400 [=====] - 2s 244us/sample - loss: 0.1055 - acc:
0.9695
8400/8400 [=====] - 2s 244us/sample - loss: 0.1092 - acc:
0.9654
8400/8400 [=====] - 2s 245us/sample - loss: 0.1017 - acc:
0.9698
8400/8400 [=====] - 2s 229us/sample - loss: 0.1103 - acc:
0.9656
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: adam , 5
epochs. Accuracy: 0.967976188659668
8400/8400 [=====] - 2s 234us/sample - loss: 2.3026 - acc:
0.0992
8400/8400 [=====] - 2s 241us/sample - loss: 2.3026 - acc:
0.0915
8400/8400 [=====] - 2s 235us/sample - loss: 2.3033 - acc:
0.1019
8400/8400 [=====] - 2s 250us/sample - loss: 2.3026 - acc:
0.0979
8400/8400 [=====] - 2s 246us/sample - loss: 2.3026 - acc:
0.1014
Hidden size 125 , Activation 1: softmax , Activation 2: relu , Optimizer: adam , 5
epochs. Accuracy: 0.09838095158338547
8400/8400 [=====] - 2s 257us/sample - loss: 0.4337 - acc:
0.8682
8400/8400 [=====] - 2s 250us/sample - loss: 0.4140 - acc:
0.8846
8400/8400 [=====] - 2s 260us/sample - loss: 0.5472 - acc:
0.8146
8400/8400 [=====] - 2s 251us/sample - loss: 0.5481 - acc:
0.8570
8400/8400 [=====] - 2s 252us/sample - loss: 0.5995 - acc:
0.7971
Hidden size 125 , Activation 1: softmax , Activation 2: softmax , Optimizer: adam ,
5 epochs. Accuracy: 0.8443333387374878
```

Next, the number of epochs and the gradient descent method is optimized.

```
In [87]: for l in ['RMSprop', 'adam', 'Nadam']:
         for m in [3,5,10]:
             evalmodel(125,tf.nn.relu,tf.nn.softmax,l,m)
```

```
8400/8400 [=====] - 2s 292us/sample - loss: 0.1324 - acc:
0.9645
8400/8400 [=====] - 2s 280us/sample - loss: 0.1167 - acc:
0.9663
8400/8400 [=====] - 2s 289us/sample - loss: 0.1363 - acc:
0.9607
8400/8400 [=====] - 2s 293us/sample - loss: 0.1377 - acc:
0.9587
8400/8400 [=====] - 3s 306us/sample - loss: 0.1122 - acc:
0.9685
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: RMSprop ,
3 epochs. Accuracy: 0.9637380957603454
8400/8400 [=====] - 3s 301us/sample - loss: 0.1294 - acc:
0.9664
8400/8400 [=====] - 3s 324us/sample - loss: 0.1108 - acc:
0.9670
8400/8400 [=====] - 3s 314us/sample - loss: 0.1249 - acc:
0.9667
8400/8400 [=====] - 2s 293us/sample - loss: 0.1061 - acc:
0.9708
8400/8400 [=====] - 3s 312us/sample - loss: 0.1035 - acc:
0.9707
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: RMSprop ,
5 epochs. Accuracy: 0.9683333277702332
8400/8400 [=====] - 3s 313us/sample - loss: 0.1285 - acc:
0.9707
8400/8400 [=====] - 2s 295us/sample - loss: 0.1060 - acc:
0.9725
8400/8400 [=====] - 3s 317us/sample - loss: 0.1215 - acc:
0.9688
8400/8400 [=====] - 3s 314us/sample - loss: 0.1188 - acc:
0.9702
8400/8400 [=====] - 3s 317us/sample - loss: 0.1115 - acc:
0.9724
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: RMSprop ,
10 epochs. Accuracy: 0.9709285736083985
8400/8400 [=====] - 3s 308us/sample - loss: 0.1288 - acc:
0.9629
8400/8400 [=====] - 3s 323us/sample - loss: 0.1069 - acc:
0.9682
8400/8400 [=====] - 3s 311us/sample - loss: 0.1240 - acc:
0.9642
8400/8400 [=====] - 3s 323us/sample - loss: 0.1203 - acc:
0.9640
8400/8400 [=====] - 3s 322us/sample - loss: 0.1105 - acc:
0.9670
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: adam , 3
epochs. Accuracy: 0.9652618885040283
8400/8400 [=====] - 3s 321us/sample - loss: 0.1033 - acc:
0.9701
8400/8400 [=====] - 3s 321us/sample - loss: 0.0991 - acc:
0.9700
8400/8400 [=====] - 3s 325us/sample - loss: 0.1202 - acc:
```

0.9657
8400/8400 [=====] - 3s 324us/sample - loss: 0.0986 - acc:
0.9705
8400/8400 [=====] - 3s 326us/sample - loss: 0.0935 - acc:
0.9715
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: adam , 5
epochs. Accuracy: 0.9695714354515076
8400/8400 [=====] - 3s 320us/sample - loss: 0.1099 - acc:
0.9700
8400/8400 [=====] - 3s 341us/sample - loss: 0.1178 - acc:
0.9698
8400/8400 [=====] - 3s 342us/sample - loss: 0.1233 - acc:
0.9682
8400/8400 [=====] - 3s 333us/sample - loss: 0.0969 - acc:
0.9724
8400/8400 [=====] - 3s 336us/sample - loss: 0.1007 - acc:
0.9730
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: adam , 10
epochs. Accuracy: 0.9706666707992554
8400/8400 [=====] - 4s 454us/sample - loss: 0.1203 - acc:
0.9655
8400/8400 [=====] - 3s 338us/sample - loss: 0.1204 - acc:
0.9635
8400/8400 [=====] - 3s 356us/sample - loss: 0.1315 - acc:
0.9592
8400/8400 [=====] - 3s 348us/sample - loss: 0.1260 - acc:
0.9614
8400/8400 [=====] - 3s 359us/sample - loss: 0.1231 - acc:
0.9610
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 3
epochs. Accuracy: 0.9620952486991883
8400/8400 [=====] - 3s 358us/sample - loss: 0.0941 - acc:
0.9727
8400/8400 [=====] - 3s 359us/sample - loss: 0.1035 - acc:
0.9700
8400/8400 [=====] - 3s 359us/sample - loss: 0.1112 - acc:
0.9679
8400/8400 [=====] - 3s 349us/sample - loss: 0.1098 - acc:
0.9660
8400/8400 [=====] - 3s 403us/sample - loss: 0.1094 - acc:
0.9671
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 5
epochs. Accuracy: 0.9687381029129029
8400/8400 [=====] - 3s 375us/sample - loss: 0.1055 - acc:
0.9725
8400/8400 [=====] - 3s 407us/sample - loss: 0.1067 - acc:
0.9724
8400/8400 [=====] - 3s 372us/sample - loss: 0.1164 - acc:
0.9713
8400/8400 [=====] - 3s 392us/sample - loss: 0.1080 - acc:
0.9701
8400/8400 [=====] - 3s 372us/sample - loss: 0.1070 - acc:
0.9707
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 1
0 epochs. Accuracy: 0.971404767036438

Finally, the size of the hidden layer is optimized.

```
In [90]: for i in [125, 250, 400]:
          evalmodel(i,tf.nn.relu,tf.nn.softmax,'Nadam',10)

8400/8400 [=====] - 4s 420us/sample - loss: 0.1130 - acc:
0.9721
8400/8400 [=====] - 3s 397us/sample - loss: 0.1005 - acc:
0.9735
8400/8400 [=====] - 3s 405us/sample - loss: 0.1179 - acc:
0.9710
8400/8400 [=====] - 3s 390us/sample - loss: 0.1179 - acc:
0.9686
8400/8400 [=====] - 3s 388us/sample - loss: 0.0960 - acc:
0.9732
Hidden size 125 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 1
0 epochs. Accuracy: 0.971666669845581
8400/8400 [=====] - 3s 407us/sample - loss: 0.1057 - acc:
0.9719
8400/8400 [=====] - 3s 406us/sample - loss: 0.0970 - acc:
0.9760
8400/8400 [=====] - 4s 421us/sample - loss: 0.1108 - acc:
0.9752
8400/8400 [=====] - 4s 429us/sample - loss: 0.1229 - acc:
0.9661
8400/8400 [=====] - 4s 443us/sample - loss: 0.1128 - acc:
0.9729
Hidden size 250 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 1
0 epochs. Accuracy: 0.9724047541618347
8400/8400 [=====] - 4s 454us/sample - loss: 0.1277 - acc:
0.9726
8400/8400 [=====] - 4s 453us/sample - loss: 0.0951 - acc:
0.9765
8400/8400 [=====] - 4s 452us/sample - loss: 0.1243 - acc:
0.9713
8400/8400 [=====] - 4s 455us/sample - loss: 0.1072 - acc:
0.9735
8400/8400 [=====] - 4s 459us/sample - loss: 0.1237 - acc:
0.9707
Hidden size 400 , Activation 1: relu , Activation 2: softmax , Optimizer: Nadam , 1
0 epochs. Accuracy: 0.9729285717010498
```

A similar model with a similar number of total weights but using 2 hidden layers instead of 1 is evaluated below.


```
In [94]: meanacc = 0
         for i in range(1,6):
             CVsets = CVsplit(i)
             model = keras.Sequential([keras.layers.Dense(300, activation=tf.nn.relu),keras.l
                 keras.layers.Dense(10, activation = tf.nn.softmax)])
             model.compile(optimizer='Nadam', loss='sparse_categorical_crossentropy', metrics
             model.fit(np.array(CVsets[1].iloc[:,1:]), np.array(CVsets[1].label), epochs=10,
             resultloss, resultacc = model.evaluate(np.array(CVsets[0].iloc[:,1:]), np.array(
             meanacc += resultacc
         meanacc /=5
         print('Accuracy: ',meanacc)
```

```
8400/8400 [=====] - 4s 488us/sample - loss: 0.1281 - acc:
0.9723
8400/8400 [=====] - 4s 489us/sample - loss: 0.1091 - acc:
0.9782
8400/8400 [=====] - 4s 500us/sample - loss: 0.1417 - acc:
0.9726
8400/8400 [=====] - 4s 497us/sample - loss: 0.1147 - acc:
0.9745
8400/8400 [=====] - 4s 488us/sample - loss: 0.1124 - acc:
0.9762
Accuracy:  0.9747618913650513
```

The above model was the best performing model thus far, in terms of cross validation. The model will be fit and predictions made on the test set.

```
In [95]: model.fit(np.array(train.iloc[:,1:]), np.array(train.label), epochs=10, verbose = 0)
         resultloss, resultacc = model.evaluate(np.array(train.iloc[:,1:]), np.array(train.la
         print('Accuracy: ',resultacc)
```

```
42000/42000 [=====] - 12s 288us/sample - loss: 0.0100 - ac
c: 0.9971
Accuracy:  0.99707144
```

```
In [96]: predictions = model.predict(np.array(test))
```

```
In [102]: pd.DataFrame(np.argmax(predictions,axis=1)).to_csv(r'C:\Datasets\digitspredict.csv')
```