

Store Sales Prediction

```
In [27]: library(ggplot2)
library(caret)
library(gridExtra)
library(dplyr)
library(stringr)

...
```

Data Exploration and Cleaning

```
In [5]: df <- read.csv('C:/Datasets/SalesTrain.txt',na.strings='')
head(df,5)
length(df[,1])
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	O
FDA15	9.30	Low Fat	0.01604730	Dairy	249.8092	OUT049	
DRC01	5.92	Regular	0.01927822	Soft Drinks	48.2692	OUT018	
FDN15	17.50	Low Fat	0.01676007	Meat	141.6180	OUT049	
FDX07	19.20	Regular	0.00000000	Fruits and Vegetables	182.0950	OUT010	
NCD19	8.93	Low Fat	0.00000000	Household	53.8614	OUT013	

8523

Check for duplicates

```
In [6]: sum(duplicated(df))

0
```

Summarize Data

In [7]: summary(df)

```
Item_Identifier  Item_Weight      Item_Fat_Content  Item_Visibility
FDG33   :   10   Min.      : 4.555    LF      : 316    Min.      :0.00000
FDW13   :   10   1st Qu.: 8.774    low fat: 112    1st Qu.:0.02699
DRE49   :    9   Median :12.600    Low Fat:5089   Median :0.05393
DRN47   :    9   Mean    :12.858    reg      : 117   Mean    :0.06613
FDD38   :    9   3rd Qu.:16.850    Regular:2889   3rd Qu.:0.09459
FDF52   :    9   Max.     :21.350                Max.     :0.32839
(Other):8467   NA's     :1463

      Item_Type      Item_MRP      Outlet_Identifier
Fruits and Vegetables:1232  Min.      : 31.29  OUT027 : 935
Snack Foods           :1200  1st Qu.: 93.83  OUT013 : 932
Household             : 910  Median :143.01  OUT035 : 930
Frozen Foods          : 856  Mean    :140.99  OUT046 : 930
Dairy                 : 682  3rd Qu.:185.64  OUT049 : 930
Canned                : 649  Max.     :266.89  OUT045 : 929
(Other)               :2994                (Other):2937
Outlet_Establishment_Year  Outlet_Size  Outlet_Location_Type
Min.      :1985                High   : 932  Tier 1:2388
1st Qu.:1987                Medium:2793  Tier 2:2785
Median :1999                Small  :2388  Tier 3:3350
Mean     :1998                NA's   :2410
3rd Qu.:2004
Max.     :2009

      Outlet_Type      Item_Outlet_Sales
Grocery Store      :1083  Min.      :   33.29
Supermarket Type1:5577  1st Qu.:   834.25
Supermarket Type2: 928  Median : 1794.33
Supermarket Type3: 935  Mean     : 2181.29
                        3rd Qu.: 3101.30
                        Max.     :13086.97
```

In [8]: names(df)[1]
length(unique(df[,1]))
names(df)[5]
length(unique(df[,5]))
names(df)[7]
length(unique(df[,7]))

'Item_Identifier'

1559

'Item_Type'

16

'Outlet_Identifier'

10

Item identifier has many levels and likely cannot be included as a feature as is. Prefixes are explored.

```
In [9]: itemprefix <- unique(substring(df[,1],1,3))
paste(length(itemprefix), 'levels')
itemprefix

itemprefix2 <- unique(substring(df[,1],1,2))
itemprefix2
```

'71 levels'

'FDA' 'DRC' 'FDN' 'FDX' 'NCD' 'FDP' 'FDO' 'FDH' 'FDU' 'FDY' 'FDS' 'FDF' 'NCB'
'DRI' 'FDW' 'FDC' 'FDR' 'FDV' 'DRJ' 'FDE' 'NCS' 'DRH' 'NCX' 'DRZ' 'FDB' 'FDK'
'FDL' 'FDM' 'NCP' 'NCL' 'DRK' 'FDI' 'FDZ' 'NCI' 'FDJ' 'FDG' 'NCZ' 'FDQ' 'FDD'
'DRG' 'NCR' 'FDT' 'DRB' 'DRE' 'DRA' 'NCF' 'NCH' 'NCO' 'NCN' 'NCC' 'DRD' 'DRF'
'DRL' 'NCM' 'NCU' 'DRY' 'NCW' 'DRM' 'NCT' 'NCQ' 'DRP' 'DRQ' 'NCK' 'NCY' 'DRN'
'NCA' 'NCE' 'NCJ' 'NCV' 'NCG' 'DRO'

'FD' 'DR' 'NC'

2 letter prefix only has 3 levels. 3 level prefix has 71 levels. Overfitting will be evaluated before deciding which prefix to keep in the model. For now both will be added to the data frame.

```
In [10]: df$ItemPrefix2 <- factor(substring(df[,1],1,2))
df$ItemPrefix3 <- factor(substring(df[,1],1,3))
df$ItemNumber <- as.numeric(substring(df[,1],4,5))
df<-df[, -1]
head(df,5)
```

Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishm
9.30	Low Fat	0.01604730	Dairy	249.8092	OUT049	
5.92	Regular	0.01927822	Soft Drinks	48.2692	OUT018	
17.50	Low Fat	0.01676007	Meat	141.6180	OUT049	
19.20	Regular	0.00000000	Fruits and Vegetables	182.0950	OUT010	
8.93	Low Fat	0.00000000	Household	53.8614	OUT013	

The fat content should only have 2 levels. This is corrected below.

```
In [11]: df$Item_Fat_Content <- factor(ifelse(df$Item_Fat_Content %in% c('reg','Regular'), 'r
head(df,5)
```

Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishm
9.30	low	0.01604730	Dairy	249.8092	OUT049	
5.92	reg	0.01927822	Soft Drinks	48.2692	OUT018	
17.50	low	0.01676007	Meat	141.6180	OUT049	
19.20	reg	0.00000000	Fruits and Vegetables	182.0950	OUT010	
8.93	low	0.00000000	Household	53.8614	OUT013	

Summary of data after cleaning

In [12]: summary(df)

```
Item_Weight      Item_Fat_Content Item_Visibility
Min.   : 4.555    low:5517          Min.      :0.00000
1st Qu.: 8.774    reg:3006          1st Qu.:0.02699
Median :12.600                    Median :0.05393
Mean   :12.858                    Mean   :0.06613
3rd Qu.:16.850                    3rd Qu.:0.09459
Max.    :21.350                    Max.    :0.32839
NA's     :1463

      Item_Type      Item_MRP      Outlet_Identifier
Fruits and Vegetables:1232  Min.   : 31.29  OUT027 : 935
Snack Foods           :1200  1st Qu.: 93.83  OUT013 : 932
Household             : 910  Median :143.01  OUT035 : 930
Frozen Foods          : 856  Mean    :140.99  OUT046 : 930
Dairy                 : 682  3rd Qu.:185.64  OUT049 : 930
Canned                : 649  Max.    :266.89  OUT045 : 929
(Other)               :2994                (Other):2937
Outlet_Establishment_Year Outlet_Size  Outlet_Location_Type
Min.   :1985                High   : 932  Tier 1:2388
1st Qu.:1987                Medium:2793  Tier 2:2785
Median :1999                Small  :2388  Tier 3:3350
Mean   :1998                NA's   :2410
3rd Qu.:2004
Max.    :2009

      Outlet_Type  Item_Outlet_Sales  ItemPrefix2  ItemPrefix3
Grocery Store    :1083  Min.   : 33.29  DR: 799    FDX    : 295
Supermarket Type1:5577  1st Qu.: 834.25  FD:6125    FDT    : 281
Supermarket Type2: 928  Median : 1794.33  NC:1599    FDS    : 279
Supermarket Type3: 935  Mean    : 2181.29                FDR    : 275
                        3rd Qu.: 3101.30                FDW    : 271
                        Max.    :13086.97                FDY    : 271
                                                (Other):6851

      ItemNumber
Min.   : 1.00
1st Qu.:15.00
Median :30.00
Mean   :30.33
3rd Qu.:45.00
Max.    :60.00
```

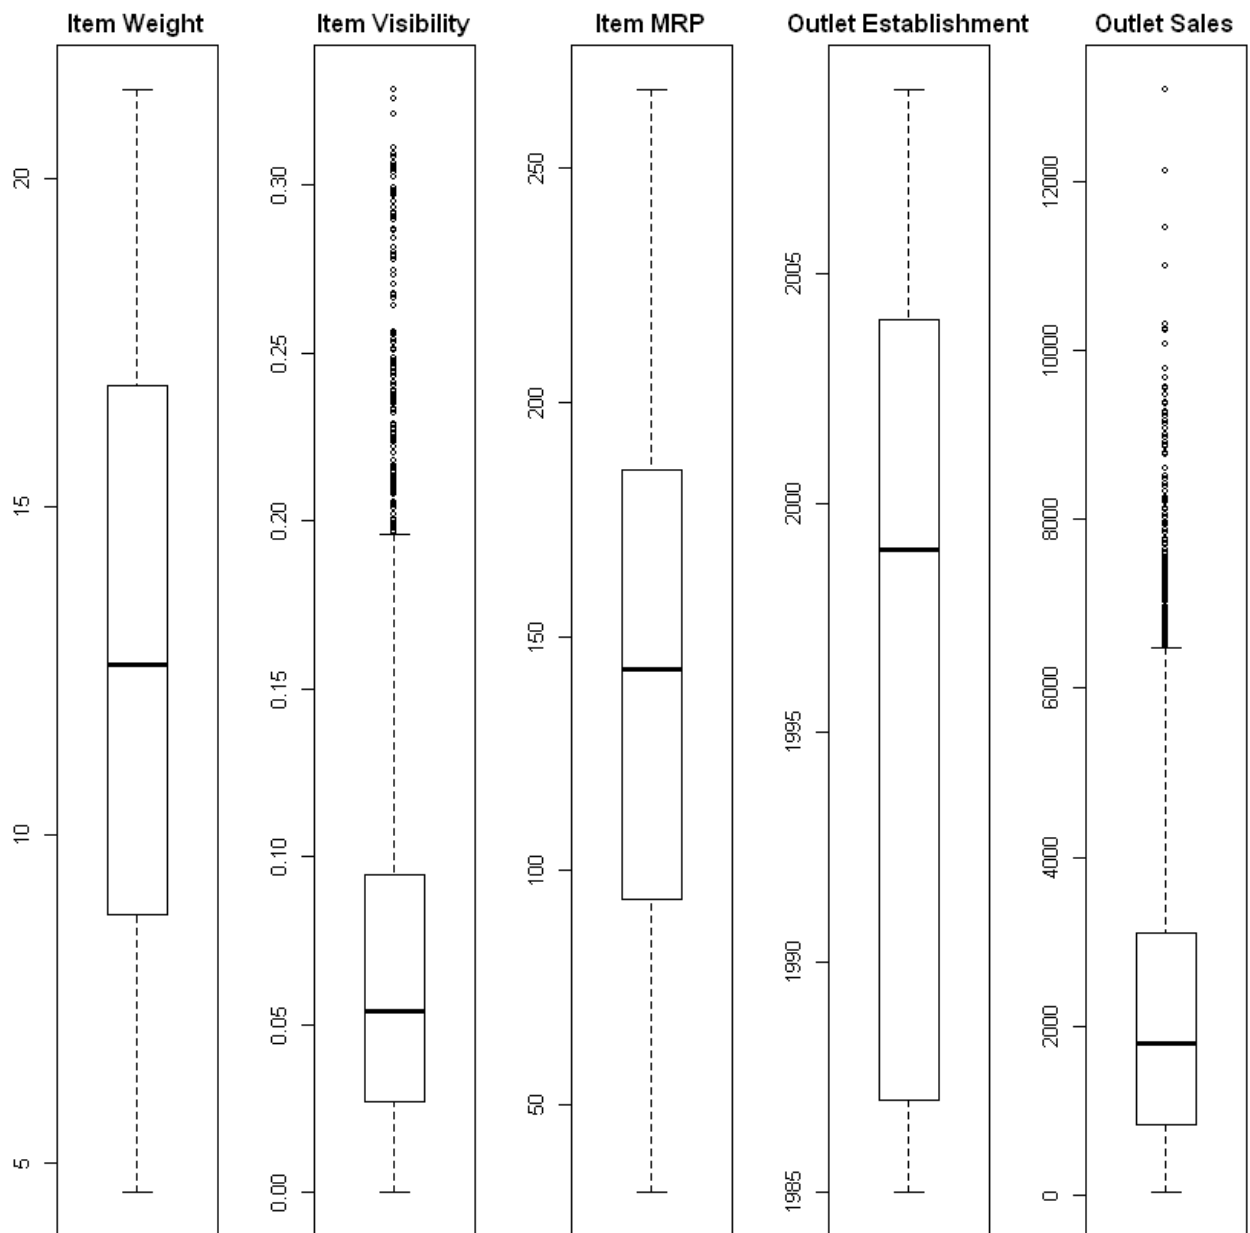
Check for near zero variance.

```
In [13]: nearZeroVar(df, saveMetrics=TRUE)
```

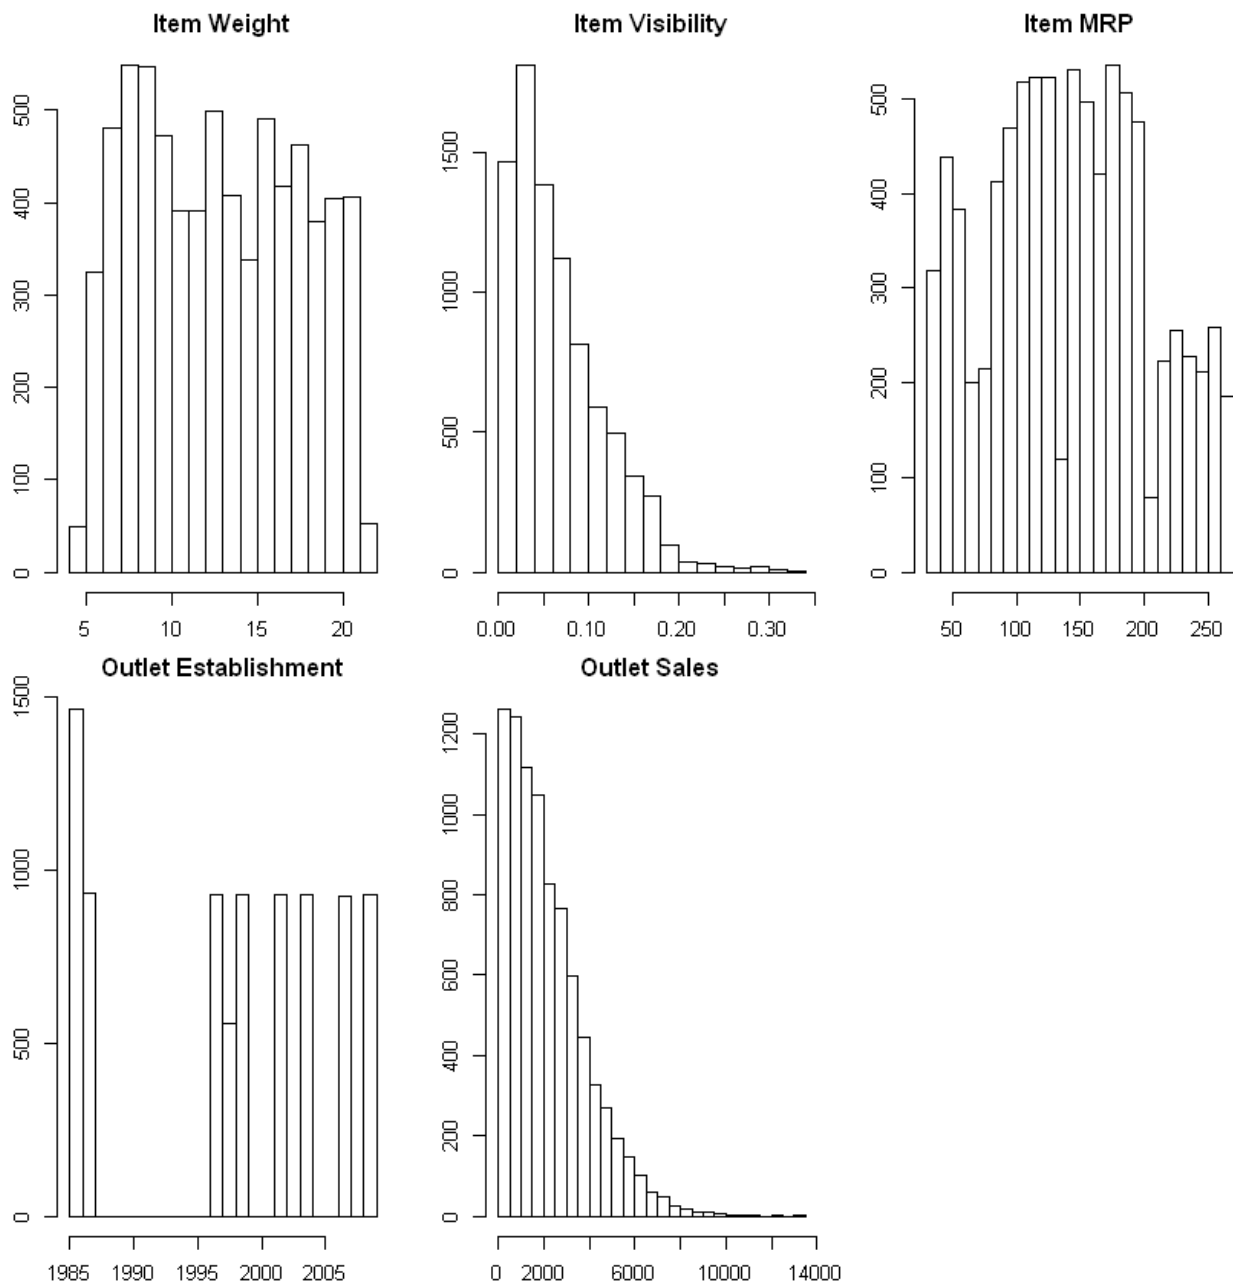
	freqRatio	percentUnique	zeroVar	nzv
Item_Weight	1.048780	4.86917752	FALSE	FALSE
Item_Fat_Content	1.835329	0.02346592	FALSE	FALSE
Item_Visibility	175.333333	92.45570808	FALSE	FALSE
Item_Type	1.026667	0.18772733	FALSE	FALSE
Item_MRP	1.166667	69.67030388	FALSE	FALSE
Outlet_Identifier	1.003219	0.11732958	FALSE	FALSE
Outlet_Establishment_Year	1.569742	0.10559662	FALSE	FALSE
Outlet_Size	1.169598	0.03519887	FALSE	FALSE
Outlet_Location_Type	1.202873	0.03519887	FALSE	FALSE
Outlet_Type	5.149584	0.04693183	FALSE	FALSE
Item_Outlet_Sales	1.062500	40.98322187	FALSE	FALSE
ItemPrefix2	3.830519	0.03519887	FALSE	FALSE
ItemPrefix3	1.049822	0.83304001	FALSE	FALSE
ItemNumber	1.006410	0.70397747	FALSE	FALSE

Check for outliers

```
In [14]: par(mfrow=c(1,5), mar=c(2,2,2,2))
boxplot(df$Item_Weight,main='Item Weight')
boxplot(df$Item_Visibility,main='Item Visibility')
boxplot(df$Item_MRP,main='Item MRP')
boxplot(df$Outlet_Establishment_Year,main='Outlet Establishment')
boxplot(df$Item_Outlet_Sales,main='Outlet Sales')
```



```
In [15]: par(mfrow=c(2,3), mar=c(2,2,2,2))
hist(df$Item_Weight,main='Item Weight',breaks=20)
hist(df$Item_Visibility,main='Item Visibility',breaks=20)
hist(df$Item_MRP,main='Item MRP',breaks=20)
hist(df$Outlet_Establishment_Year,main='Outlet Establishment',breaks=20)
hist(df$Item_Outlet_Sales,main='Outlet Sales',breaks=20)
```



The two variables with extreme values are item visibility and outlet sales. However, the extreme values are numerous and are consistent with the skewed distributions observed in the histogram. Therefore they are not considered outliers.

Next, check for correlations.


```
In [16]: cor(df[,c(1,3,5,7,14)],use='complete.obs')
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	ItemNu
Item_Weight	1.00000000	-0.014047726	0.027141154	-0.011588290	0.0252
Item_Visibility	-0.01404773	1.000000000	-0.006061148	-0.016935201	-0.0265
Item_MRP	0.02714115	-0.006061148	1.000000000	-0.001656520	-0.0188
Outlet_Establishment_Year	-0.01158829	-0.016935201	-0.001656520	1.000000000	0.0011
ItemNumber	0.02528662	-0.026579321	-0.018875820	0.001171761	1.0000

```
In [17]: chisqfun <- function(a,b){  
  if(a==b) return(0)  
  else return(round(chisq.test(table(df[,a],df[,b]))$p.value,2))  
}  
matrix(mapply(chisqfun,rep(c(2,4,6,8:10,12:13),8),rep(c(2,4,6,8:10,12:13),rep(8,8)))
```

```
Warning message in chisq.test(table(df[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):  
"Chi-squared approximation may be incorrect"Warning message in chisq.test(table(d  
f[, a], df[, b])):
```

Some of the factors are not independent. The strength of correlation will be evaluated after dummy coding.

Next, factors will be coded and pre-processing will be performed.

```
In [18]: X <- predict(dummyVars(Item_Outlet_Sales~.,df,fullRank=TRUE),df)
dimnames(X)[[2]] <- gsub(' ', '', dimnames(X)[[2]])
head(X,5)
```

	Item_Weight	Item_Fat_Content.reg	Item_Visibility	Item_Type.Breads	Item_Type.Breakfast	Item_Type.C
1	9.30	0	0.01604730	0	0	
2	5.92	1	0.01927822	0	0	
3	17.50	0	0.01676007	0	0	
4	19.20	1	0.00000000	0	0	
5	8.93	0	0.00000000	0	0	

```
In [19]: preprocessparams <- preProcess(X,method=c('center','scale','bagImpute'))
processedX <- predict(preprocessparams, X)
head(processedX,5)
```

	Item_Weight	Item_Fat_Content.reg	Item_Visibility	Item_Type.Breads	Item_Type.Breakfast	Item_Type.C
1	-0.7661631	-0.7381039	-0.9706752	-0.1741831	-0.1143393	-0.2
2	-1.4940692	1.3546638	-0.9080580	-0.1741831	-0.1143393	-0.2
3	0.9997627	-0.7381039	-0.9568612	-0.1741831	-0.1143393	-0.2
4	1.3658693	1.3546638	-1.2816826	-0.1741831	-0.1143393	-0.2
5	-0.8458452	-0.7381039	-1.2816826	-0.1741831	-0.1143393	-0.2

Check for and remove any correlated dummy variables

```
In [20]: cormatrix <- cor(X,use='complete.obs')
cormatrix[is.na(cormatrix)] <- 0
findCorrelation(cormatrix,0.9,names=TRUE, exact=FALSE)
findCorrelation(cormatrix,0.9, exact=FALSE)
```

Warning message in cor(X, use = "complete.obs"):
"the standard deviation is zero"

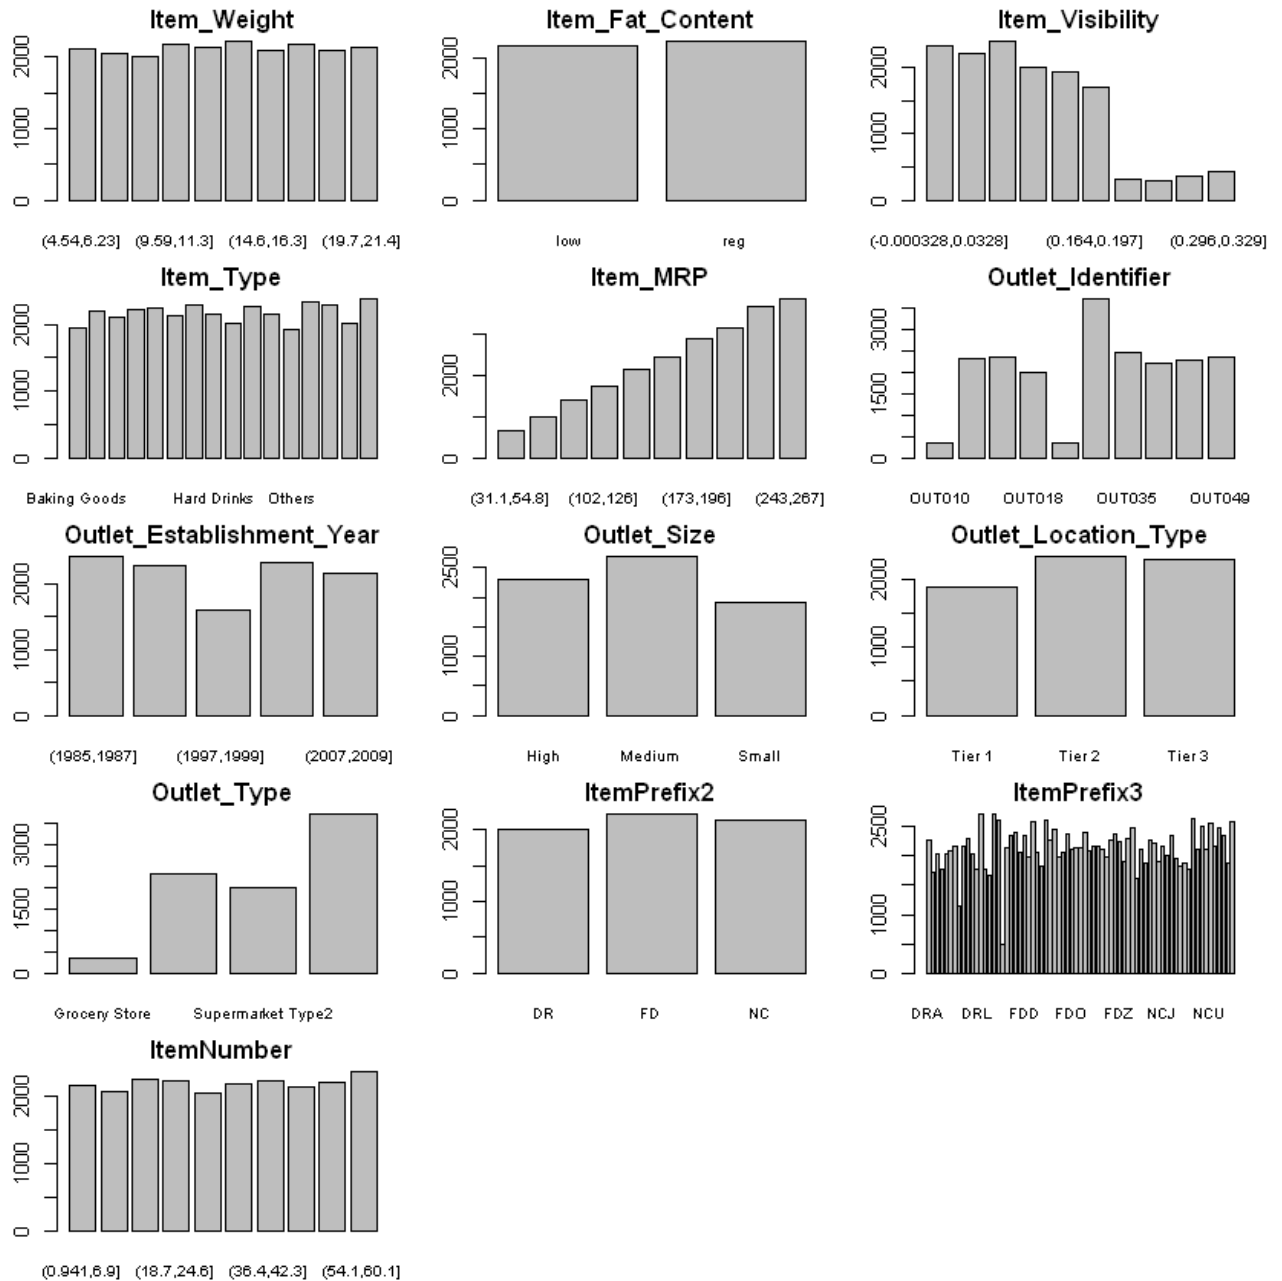
'Outlet_Identifier.OUT035' 'Outlet_Identifier.OUT018' 'Outlet_Type.SupermarketType1'

25 22 34

```
In [21]: processedX <- processedX[, -c(25,22,34)]
```

Data Visualization

```
In [22]: par(mfrow=c(5,3),mar=c(2,2,2))
for(i in 1:14){
  if(i %in% c(2,4,6,8:10,12:13)) heightmatrix <- as.matrix(aggregate(df[,11],list(
  else if(i != 11) heightmatrix <- as.matrix(aggregate(df[,11],list(cut(df[,i],10)
  if(i!=11){
    heights <- as.numeric(heightmatrix[,2])
    names(heights) <- heightmatrix[,1]
    barplot(heights, beside=TRUE, main=names(df)[i],cex.names=0.75)
  }
}
```



Item visibility, MRP, outlet identifier, and outlet type seem to have the largest effects on sales.

Linear Model

First model with all the current features.

```
In [23]: linearmod <- lm(df[,11]~processedX)
lmsummary <- summary(linearmod)
lmsummary
```

```
Call:
lm(formula = df[, 11] ~ processedX)
```


```
Residuals:
    Min       1Q   Median       3Q      Max
-4372.4  -682.0   -87.0   571.3  7880.6
```

Coefficients: (8 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2283.7186	12.9386	176.505	< 2e-16
processedXItem_Weight	-2.3546	13.8642	-0.170	0.8651
processedXItem_Fat_Content.reg	24.2491	13.8339	1.753	0.0797
processedXItem_Visibility	-15.8401	13.0961	-1.210	0.2265
processedXItem_Type.Breads	0.1974	14.4324	0.014	0.9891
processedXItem_Type.Breakfast	9.1583	13.5397	0.676	0.4988
processedXItem_Type.Canned	14.2974	16.8406	0.849	0.3959
processedXItem_Type.Dairy	-11.7418	18.2689	-0.643	0.5204
processedXItem_Type.FrozenFoods	-2.3343	17.9588	-0.130	0.8966

Significance of features is summarized below.

```
In [24]: barplot(lmsummary$coefficients[abs(lmsummary$coefficients[,3])>=1,3],main='t-values'
```



```
In [25]: sigfeatures <- lmsummary$coefficients[abs(lmsummary$coefficients[,4])<=0.05,]
sigfeatures
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2283.71861	12.93858	176.504568	0.000000e+00
processedXItem_MRP	970.89650	12.63025	76.870748	0.000000e+00
processedXOutlet_Identifier.OUT013	750.78317	33.93937	22.121307	1.900360e-105
processedXOutlet_Identifier.OUT017	102.15748	14.27344	7.157173	8.926718e-13
processedXOutlet_Identifier.OUT019	188.08163	29.37974	6.401745	1.618243e-10
processedXOutlet_Identifier.OUT027	1649.66946	38.07139	43.330953	0.000000e+00
processedXOutlet_Identifier.OUT045	278.83950	14.50041	19.229768	1.099017e-80
processedXOutlet_Identifier.OUT046	278.17601	20.44095	13.608759	9.792648e-42
processedXOutlet_Identifier.OUT049	578.17295	18.15598	31.844759	3.124907e-210
processedXOutlet_Establishment_Year	1239.59742	46.90286	26.429038	6.045229e-148
processedXOutlet_Size.Medium	-578.03530	23.82113	-24.265652	8.720852e-126
processedXItemPrefix3.NCH	-37.35915	17.17607	-2.175070	2.965258e-02
processedXItemPrefix3.NCI	-36.40136	17.91378	-2.032031	4.218181e-02
processedXItemNumber	53.04357	12.53284	4.232366	2.336990e-05

Add interactions of significant features to design matrix.

```
In [28]: sigfeaturelist <- str_replace_all(dimnames(sigfeatures)[[1]], 'processedX', '')[-1]
sigfeaturelist
```

```
'Item_MRP' 'Outlet_Identifier.OUT013' 'Outlet_Identifier.OUT017' 'Outlet_Identifier.OUT019'
'Outlet_Identifier.OUT027' 'Outlet_Identifier.OUT045' 'Outlet_Identifier.OUT046'
'Outlet_Identifier.OUT049' 'Outlet_Establishment_Year' 'Outlet_Size.Medium' 'ItemPrefix3.NCH'
'ItemPrefix3.NCI' 'ItemNumber'
```

```
In [29]: InterX <- as.data.frame(processedX)
for(a in 1:12){
  for(b in (a+1):13){
    intername <- paste(sigfeaturelist[a],sigfeaturelist[b],sep='x')
    InterX <- mutate(InterX,!!intername := processedX[,sigfeaturelist[a]]*proces
  }
}
InterX <- as.matrix(InterX)
```

Check for and remove interactions that are highly correlated with other features.

```
In [30]: intermatrix <-cor(InterX,use='complete.obs')
intermatrix[is.na(cormatrix)] <- 0
diag(intercormatrix) <- 0
maxcor <- apply(intercormatrix,1,max)
filteredindex <- 1:length(maxcor)
filteredindex <- filteredindex[(maxcor >= 0.9)&(filteredindex>98)]
filteredindex
```

```
125 126 135 144 145 153 165 170 171
```

```
In [31]: InterX <-InterX[,-filteredindex]
```

Run linear model with interactions.

```
In [32]: linearmodinter <-lm(df[,11]~InterX)
lmintersummary <- summary(linearmodinter)
lmintersummary
```

Call:

```
lm(formula = df[, 11] ~ InterX)
```

Residuals:

Min	1Q	Median	3Q	Max
-5514.8	-534.7	-53.6	452.4	6873.6

Coefficients: (36 not defined because of singularities)

	Estimate	Std. Error
(Intercept)	2284.92041	12.28974
InterXItem_Weight	-6.06602	13.16625
InterXItem_Fat_Content.reg	17.44366	13.12486
InterXItem_Visibility	-13.72552	12.43355
InterXItem_Type.Breads	1.64473	13.68194
InterXItem_Type.Breakfast	8.65375	12.83275
InterXItem_Type.Canned	10.90327	15.97262
InterXItem_Type.Dairy	-9.23726	17.32069
InterXItem_Type.FrozenFoods	-6.18434	17.02545
InterXItem_Type.FruitVegetables	8.58571	18.45426

Filter out insignificant features.

```
In [33]: siginter <-lmintersummary$coefficients[lmlintersummary$coefficients[,4]<=0.05,]
siginterlist <- str_replace_all(dimnames(siginter)[[1]], 'InterX', '')[-1]
siginterlist <- siginterlist[13:25]
siginterlist
```

```
'Item_MRPxOutlet_Identifier.OUT013' 'Item_MRPxOutlet_Identifier.OUT017'
'Item_MRPxOutlet_Identifier.OUT019' 'Item_MRPxOutlet_Identifier.OUT027'
'Item_MRPxOutlet_Identifier.OUT045' 'Item_MRPxOutlet_Identifier.OUT046'
'Item_MRPxOutlet_Identifier.OUT049' 'Item_MRPxOutlet_Establishment_Year'
'Item_MRPxOutlet_Size.Medium' 'Item_MRPxItemPrefix3.NCH' 'Item_MRPxItemPrefix3.NCI'
'Item_MRPxItemNumber' 'Outlet_Identifier.OUT013xItemPrefix3.NCH'
```

```
In [34]: fullX <- as.data.frame(processedX)
  for(a in siginterlist){
    fullX <- mutate(fullX, !!a := InterX[,a])
  }
  fullX <- as.matrix(fullX)

  linearmodfull <- lm(df[,11]~fullX)
  lmfullsummary <- summary(linearmodfull)
  lmfullsummary
```

Call:
lm(formula = df[, 11] ~ fullX)

Residuals:

Min	1Q	Median	3Q	Max
-5548.3	-539.1	-51.3	455.9	6917.1

Coefficients: (8 not defined because of singularities)

	Estimate	Std. Error	t value
(Intercept)	2285.3341	12.2749	186.179
fullXItem_Weight	-6.3010	13.1433	-0.479
fullXItem_Fat_Content.reg	17.2192	13.1198	1.312
fullXItem_Visibility	-12.4791	12.4140	-1.005
fullXItem_Type.Breads	1.6217	13.6760	0.119
fullXItem_Type.Breakfast	8.9046	12.8283	0.694
fullXItem_Type.Canned	11.1434	15.9631	0.698
fullXItem_Type.Dairy	-9.5597	17.3132	-0.552
fullXItem_Type.FrozenFoods	-6.6393	17.0187	-0.390

Add square and cube features for continuous features.


```
In [35]: nonlinX <- as.data.frame(fullX)
continuouslist <- c('Item_Weight', 'Item_MRP', 'Item_Visibility', 'Outlet_Establishment
for(a in continuouslist){
  sqname <- paste(a, 'sqr', sep='.')
  cubname <- paste(a, 'cub', sep='.')
  nonlinX <- mutate(nonlinX, !!sqname := fullX[,a]^2)
  nonlinX <- mutate(nonlinX, !!cubname := fullX[,a]^3)
}
nonlinX <- as.matrix(nonlinX)

nonlinearmod <- lm(df[,11]~nonlinX)
nlmsummary <- summary(nonlinearmod)
nlmsummary
```

Call:

```
lm(formula = df[, 11] ~ nonlinX)
```

Residuals:

Min	1Q	Median	3Q	Max
-5451.3	-533.3	-50.9	456.1	6908.8

Coefficients: (10 not defined because of singularities)

	Estimate	Std. Error	t value
(Intercept)	2309.78682	24.45734	94.441
nonlinXItem_Weight	47.52413	32.04466	1.483
nonlinXItem_Fat_Content.reg	16.73881	13.12263	1.276
nonlinXItem_Visibility	-19.48536	15.44958	-1.261
nonlinXItem_Type.Breads	0.34177	13.68739	0.025
nonlinXItem_Type.Breakfast	8.97961	12.85298	0.699
nonlinXItem_Type.Canned	10.17210	15.97353	0.637
nonlinXItem_Type.Dairy	-8.72684	17.34316	-0.503
nonlinXItem_Type.FrozenFoods	-7.47347	17.04192	-0.439

Only the cube of item weight and MRP is significant (at alpha = 0.1), so other cube and square terms are removed.

```
In [36]: finalX <- as.data.frame(fullX)
continuouslist <- c('Item_Weight', 'Item_MRP')
for(a in continuouslist){
  cubname <- paste(a, 'cub', sep='.')
  finalX <- mutate(finalX, !!cubname := fullX[,a]^3)
}
finalX <- as.matrix(finalX)
finallinearmod <- lm(df[,11]~finalX)
finalsummary <- summary(finallinearmod)
finalsummary
```

Call:

```
lm(formula = df[, 11] ~ finalX)
```

Residuals:

Min	1Q	Median	3Q	Max
-5459.4	-531.9	-53.0	452.7	6914.5

Coefficients: (8 not defined because of singularities)

	Estimate	Std. Error	t value
(Intercept)	2290.3463	12.3960	184.765
finalXItem_Weight	52.2886	31.8723	1.641
finalXItem_Fat_Content.reg	16.4154	13.1180	1.251
finalXItem_Visibility	-13.3536	12.4133	-1.076
finalXItem_Type.Breads	0.3820	13.6849	0.028
finalXItem_Type.Breakfast	8.8011	12.8506	0.685
finalXItem_Type.Canned	10.5802	15.9649	0.663
finalXItem_Type.Dairy	-8.8606	17.3183	-0.512
finalXItem_Type.FrozenFoods	-7.2372	17.0190	-0.425

The singularities are caused by features that are linear combinations of other features. These are filtered out by taking only features listed in the linear model coefficient summary.

```
In [46]: finalfeaturelist <- str_replace_all(dimnames(finallmsummary$coefficients)[[1]], 'finalX')
finalX <- finalX[,finalfeaturelist]
finallinearmod <- lm(df[,11]~finalX)
finallmsummary <- summary(finallinearmod)
finallmsummary
```

Call:
lm(formula = df[, 11] ~ finalX)

Residuals:

Min	1Q	Median	3Q	Max
-5459.4	-531.9	-53.0	452.7	6914.5

Coefficients:


	Estimate	Std. Error	t value
(Intercept)	2290.3463	12.3960	184.765
finalXItem_Weight	52.2886	31.8723	1.641
finalXItem_Fat_Content.reg	16.4154	13.1180	1.251
finalXItem_Visibility	-13.3536	12.4133	-1.076
finalXItem_Type.Breads	0.3820	13.6849	0.028
finalXItem_Type.Breakfast	8.8011	12.8506	0.685
finalXItem_Type.Canned	10.5802	15.9649	0.663
finalXItem_Type.Dairy	-8.8606	17.3183	-0.512
finalXItem_Type.FrozenFoods	-7.2372	17.0190	-0.425

Check for normality of residuals

```
In [47]: hist(finallmsummary$residuals,50, main='Histogram of Residuals')
shapiro.test(sample(finallmsummary$residuals,5000))
```

Shapiro-Wilk normality test


```
data: sample(finallmsummary$residuals, 5000)
W = 0.95835, p-value < 2.2e-16
```



The residual distribution has heavy tails, possibly indicating a large influence of features not captured in the model or in the data set.

The significance of features in the model is summarized below.

◀ [Redacted] ▶



In [49]: `finalImsummary$coefficients[finalImsummary$coefficients[,4]<=0.1,]`

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2290.34625	12.39601	184.764729	0.000000e+00
finalXItem_MRP	1066.70848	27.22300	39.184090	1.302347e-308
finalXOutlet_Identifier.OUT013	757.53434	32.14573	23.565630	5.852413e-119
finalXOutlet_Identifier.OUT017	102.78752	13.51861	7.603411	3.194981e-14
finalXOutlet_Identifier.OUT019	188.75018	27.82621	6.783181	1.255057e-11
finalXOutlet_Identifier.OUT027	1657.92788	36.05383	45.984797	0.000000e+00
finalXOutlet_Identifier.OUT045	280.17338	13.73227	20.402561	2.382812e-90
finalXOutlet_Identifier.OUT046	282.81046	19.36473	14.604412	1.007887e-47
finalXOutlet_Identifier.OUT049	580.33109	17.19359	33.752763	2.506709e-234
finalXOutlet_Establishment_Year	1247.30512	44.41542	28.082702	6.066573e-166
finalXOutlet_Size.Medium	-575.99995	22.56530	-25.525916	1.726855e-138
finalXItemPrefix3.FDT	-33.18683	16.42859	-2.020065	4.340826e-02
finalXItemPrefix3.FDW	-28.90952	16.27265	-1.776571	7.567502e-02
finalXItemPrefix3.NCC	-27.80474	15.79873	-1.759935	7.845516e-02
finalXItemPrefix3.NCI	-39.57200	16.97821	-2.330752	1.978995e-02
finalXItemPrefix3.NCT	-27.45584	15.69773	-1.749033	8.032187e-02
finalXItemNumber	42.89486	11.89430	3.606338	3.123483e-04
finalXItem_MRPxOutlet_Identifier.OUT013	358.82433	31.75561	11.299556	2.146993e-29
finalXItem_MRPxOutlet_Identifier.OUT017	43.16997	13.43739	3.212676	1.320001e-03
finalXItem_MRPxOutlet_Identifier.OUT019	85.70648	27.86869	3.075368	2.109179e-03
finalXItem_MRPxOutlet_Identifier.OUT027	718.87112	35.43004	20.289877	2.129859e-89
finalXItem_MRPxOutlet_Identifier.OUT045	119.19688	13.65513	8.729091	3.061507e-18
finalXItem_MRPxOutlet_Identifier.OUT046	106.75557	19.24694	5.546625	3.000675e-08
finalXItem_MRPxOutlet_Identifier.OUT049	260.89385	17.07128	15.282615	4.975527e-52
finalXItem_MRPxOutlet_Establishment_Year	569.37033	43.67207	13.037402	1.767222e-38
finalXItem_MRPxOutlet_Size.Medium	-253.85440	22.59335	-11.235804	4.380714e-29
finalXItem_MRPxItemPrefix3.NCH	-24.23002	13.45744	-1.800493	7.181867e-02
finalXItem_MRPxItemPrefix3.NCI	-31.60884	10.79118	-2.929137	3.408177e-03
finalXItem_MRPxItemNumber	37.53014	11.95861	3.138337	1.704971e-03
finalXItem_Weight.cub	-33.51879	16.37811	-2.046560	4.073228e-02
finalXItem_MRP.cub	-23.52734	11.49696	-2.046396	4.074845e-02

Cross validation is performed to check for overfitting.

```
In [51]: fulllinearmodel <- train(finalX,df[,11],method='lm',trControl=trainControl(method='r
fulllinearmodel
```

Linear Regression

8523 samples
113 predictor

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 3 times)

Summary of sample sizes: 6819, 6817, 6819, 6819, 6818, 6819, ...

Resampling results:

RMSE	Rsquared	MAE
1077.878	0.6014958	759.3767

Tuning parameter 'intercept' was held constant at a value of TRUE

The RMSE is comparable to that of the model fit to the entire data set, indicating overfitting is minimal.
Therefore the model will not be reduced further, nor will any regularization penalties be added.

Decision Tree

```
In [64]: Grid <- expand.grid(cp=c(0.0005,0.001,0.002,0.003,0.004,0.005))
treemodel <- train(processedX,df[,11],method='rpart',trControl=trainControl(method='
treemodel
```

CART

8523 samples
106 predictor

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 3 times)

Summary of sample sizes: 6817, 6818, 6819, 6820, 6818, 6818, ...

Resampling results across tuning parameters:

cp	RMSE	Rsquared	MAE
5e-04	1148.888	0.5520473	793.2892
1e-03	1105.266	0.5813093	773.1895
2e-03	1098.391	0.5859018	782.0021
3e-03	1101.527	0.5835262	785.5887
4e-03	1102.768	0.5825808	786.9578
5e-03	1106.873	0.5794607	790.0889

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was cp = 0.002.

```
In [69]: Grid <- expand.grid(cp=c(0.0016,0.0017,0.0018,0.0019,0.002))
treemodel <- train(processedX,df[,11],method='rpart',trControl=trainControl(method='
treemodel
```

CART

8523 samples
106 predictor

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 3 times)


Summary of sample sizes: 6818, 6818, 6818, 6819, 6819, 6817, ...

Resampling results across tuning parameters:

cp	RMSE	Rsquared	MAE
0.0016	1097.356	0.5868223	779.9323
0.0017	1097.245	0.5868887	780.0096
0.0018	1097.269	0.5868148	780.4726
0.0019	1097.308	0.5867609	780.5750
0.0020	1098.377	0.5859599	781.1463

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was cp = 0.0017.


```
In [76]: barplot(treemodel$finalModel$variable.importance,main='Feature Importance',cex.names
```



Model Selection

The linear model is slightly better than the decision tree model in terms of cross validation RMSE, therefore the linear model will be selected as the final model.

The test data set is processed the same way as the training set.

```

In [89]: test <- read.csv('C:/Datasets/SalesTest.txt',na.strings='')
test$Item_Outlet_Sales <- 0
test$ItemPrefix2 <- factor(substring(test[,1],1,2))
test$ItemPrefix3 <- factor(substring(test[,1],1,3))
test$ItemNumber <- as.numeric(substring(test[,1],4,5))
itemID <- test[,1]
test<-test[,-1]
test$Item_Fat_Content <- factor(ifelse(test$Item_Fat_Content %in% c('reg','Regular'))
testX <- predict(dummyVars(Item_Outlet_Sales~.,df,fullRank=TRUE),test)
dimnames(testX)[[2]] <- gsub(' ','',dimnames(testX)[[2]])
processedtestX <- predict(preprocessparams, testX)
processedtestX <- processedtestX[,-c(25,22,34)]
IntertestX <- as.data.frame(processedtestX)

for(a in 1:12){
  for(b in (a+1):13){
    intername <- paste(sigfeaturelist[a],sigfeaturelist[b],sep='x')
    IntertestX <- mutate(IntertestX,!!intername := processedtestX[,sigfeaturelis
  }
}
IntertestX <- as.matrix(IntertestX)
fulltestX <- as.data.frame(processedtestX)

for(a in siginterlist){
  fulltestX <- mutate(fulltestX,!!a := IntertestX[,a])
}
fulltestX <- as.matrix(fulltestX)

finaltestX <- as.data.frame(fulltestX)
for(a in continuouslist){
  cubname <- paste(a,'cub',sep='.')
  finaltestX <- mutate(finaltestX,!!cubname := fulltestX[,a]^3)
}
finaltestX <- as.matrix(finaltestX)
finaltestX <- finaltestX[,finalfeaturelist]

```

Then the model is used to predict sales for the test set.

```

In [90]: results <- predict(fulllinearmodel,finaltestX)

```

```

In [94]: resulttable <- as.data.frame(itemID)
resulttable$Outlet_Identifier <- test$Outlet_Identifier
resulttable$Item_Outlet_Sales <- results
write.csv(resulttable,'C:/Datasets/SalesTestPredict.csv')

```